

tempssm Reference Manual

Version 0.3.0

Akira S. Hirao, Shinya Baba, and Momoko Ichinokawa

2026-07-09

Summary

tempssm provides a practical R interface for state-space analysis of temperature time series. The package focuses on linear Gaussian state-space models estimated by Kalman filtering and smoothing, using the **KFAS** package as the computational backend (Helske, 2017).

Key Features

- Fits linear Gaussian state-space models to temperature time series.
- Represents temperature dynamics using interpretable latent components: long-term trend, seasonal variation, autoregressive structure, and optional exogenous effects.
- Supports arbitrary seasonal frequencies, while the current examples and validation focus primarily on monthly temperature data.
- Allows the autoregressive order to be specified by the user.
- Provides S3 methods for summaries, diagnostics, and plots.
- Includes time-series cross-validation tools for model evaluation.

Prior Art and Scope

tempssm provides a domain-focused workflow for analyzing temperature time series with linear Gaussian state-space models. It brings together model construction, component extraction, uncertainty summaries, residual diagnostics, visualization, and time-series cross-validation in a single R package interface tailored to temperature applications.

The package builds on established statistical methodology, including linear Gaussian state-space modeling, Kalman filtering, and Kalman smoothing. Model estimation is handled through the **KFAS** package, which provides a general framework for state-space models in R.

The initial implementation was adapted from the supplementary code provided by Baba (2024), accompanying Baba et al. (2024), which analyzed sea temperature trends using a linear Gaussian state-space model. The supplementary code is publicly available at:

<https://github.com/logics-of-blue/sea-temperature-trend-jogashima>

Compared with that prior implementation, **tempssm** extends the workflow into a reusable R package interface with input validation, documented S3 methods, tests, diagnostics, cross-validation utilities, and examples for broader temperature time-series analysis.

The next section describes the state-space model used by **tempssm** and clarifies how the trend, seasonal, autoregressive, and exogenous components are represented.

State-Space Model in tempssm

We consider an extended version of the Basic Structural Time Series Model (BSTSM) to describe temperature time series with explicit long-term trend, seasonal variability, and auto-regressive component. The observation equation is given by

$$y_t = \alpha_t + v_t. \quad (1)$$

where t denotes the time index, y_t is the observed temperature, α_t is the latent state, and v_t is the observation error.

The latent state is decomposed as

$$\alpha_t = \mu_t + s_t + r_t + E_t. \quad (2)$$

where μ_t is the long-term trend component, s_t is the seasonal component, r_t is a stationary autoregressive component, and E_t is the contribution of exogenous variables.

The level component follows a second-order stochastic process:

$$\mu_t = 2\mu_{t-1} - \mu_{t-2} + \zeta_t, \quad \zeta_t \sim \mathcal{N}(0, \sigma_\zeta^2). \quad (3)$$

where ζ_t is the process error of the long-term trend component.

The seasonal component is modeled with a sum-to-zero constraint:

$$s_t = - \sum_{i=t-f}^{t-1} s_i + \omega_t, \quad \omega_t \sim \mathcal{N}(0, \sigma_\omega^2). \quad (4)$$

where ω_t is the process error of the seasonal component and f denotes the seasonal frequency (for example, $f = 12$ for monthly data). This formulation makes seasonal effects identifiable and enables models to be constructed for time series with different temporal resolutions. By imposing the sum-to-zero constraint over one complete seasonal cycle, the seasonal component captures recurring deviations from the underlying long-term trend without introducing long-term drift. In models without a seasonal component, the seasonal term s_t is set to zero and omitted from the state equation.

The autoregressive component follows an l th-order autoregressive (AR) process:

$$r_t = \phi_1 r_{t-1} + \phi_2 r_{t-2} + \dots + \phi_l r_{t-l} + \tau_t, \quad \tau_t \sim \mathcal{N}(0, \sigma_\tau^2). \quad (5)$$

Here, τ_t denotes the process error of the autoregressive component. The process error terms (ζ_t , ω_t , and τ_t) are assumed to be mutually independent and normally distributed.

The exogenous component is defined as

$$E_t = \beta_1 x_{1,t} + \beta_2 x_{2,t} + \dots + \beta_m x_{m,t}. \quad (6)$$

The exogenous term E_t represents the influence of external factors. These may include large-scale climate indices, regional environmental variables, or other physically motivated predictors relevant to the observed temperature time series. Each exogenous variable enters the model linearly through a time-invariant regression coefficient ($\beta_1, \beta_2, \dots, \beta_m$), allowing the magnitude and direction of its contribution to be estimated jointly with the latent state components.

When applying the model to observational data, the total number of parameters to be estimated (k) depends on the order of the autoregressive component and the number of exogenous variables included in the model. For example, in a model with a second-order autoregressive component and no exogenous covariates, six parameters are estimated: the observation error variance, the process error variances associated with the long-term trend, seasonal, and autoregressive components, and the first- and second-order autoregressive coefficients (ϕ_1 and ϕ_2). In the general case with an l th-order autoregressive component and m exogenous variables, the total number of parameters is $k = 4 + l + m$.

The core implementation of parameter estimation procedure is based on the supplementary code provided by Baba et al. (2024). Parameter estimation is performed using a two-step optimization strategy recommended

by Helske (2017). In the first step, model parameters are estimated using the Nelder-Mead method (Nelder & Mead, 1965) with user-specified initial values. The resulting estimates are then used as initial values in a second optimization step based on the BFGS algorithm (Shanno, 1970). The main model-fitting function, `tempssm()`, returns both filtering and smoothing estimates. Unless otherwise stated, all results presented in this vignette are based on the smoothed estimates.

How to Use

Set Environment

Load the following R packages to run the examples below. If any packages are not installed, please install them as needed.

```
## Set libraries
library(tempssm)
library(forecast)

library(purrr)
library(tibble)
library(readr)
library(dplyr)
library(ggplot2)

library(patchwork)
```

Input Data Format

Input data for `tempssm` must be supplied as an R `ts` object, which represents a regularly spaced time series (see `?stats::ts` or <https://stat.ethz.ch/R-manual/R-devel/library/stats/html/ts.html>).

To support data preparation, the package includes utility functions that convert external observational data into `ts` objects (see Appendix).

Exercise I: Applying a State-Space Model to a Univariate Temperature Time Series

Objective

In Exercise I, we apply a linear Gaussian state-space model to a univariate temperature time series. This exercise introduces basic modeling without exogenous variables and examines the role of autoregressive dynamics.

Loading the Sea Surface Temperature (SST) Dataset

A sample sea surface temperature (SST) dataset is included in the package.

- **Dataset:** Monthly sea surface temperature (SST) off Yamaguchi Pref., Japan
- **Unit:** °C
- **Period:** February 2002 to December 2023

This dataset was created by aggregating daily SST data obtained from the Japan Meteorological Agency website (<https://www.jma.go.jp/jma/indexe.html>) into monthly values.

```
data(yamaguchi_sst) # load a ts object of SST off Niigata
head(yamaguchi_sst)
```

```
##           Jan      Feb      Mar      Apr      May      Jun
## 1982 14.85516 13.36500 13.55645 14.29933 17.72419 21.53000
```

```
summary(yamaguchi_sst)
```


```
##      Temp
## Min.   :11.45
## 1st Qu.:15.13
## Median :18.84
## Mean   :19.54
## 3rd Qu.:23.63
## Max.   :29.49
```

Plotting the Monthly SST Time Series

We begin by visualizing the monthly SST time series to examine its overall structure, including apparent trends, seasonal variability, and the presence or absence of missing observations.

```
plt_yamaguchi_sst <- forecast::autoplot(yamaguchi_sst) +
  labs(y = expression(Temp.~(degree*C)),
       x = "Time (year)") +
  ggtitle("Monthly SST off Yamaguchi, Japan") +
  theme_classic()

plot(plt_yamaguchi_sst)
```



Applying a Linear Gaussian State-Space Model

When a `ts` object containing temperature time-series data (here, `yamaguchi_sst`) is passed to the core function `tempssm()`, model construction and parameter estimation are performed together. The returned S3 object of class `tempssm` (here, `res_ar1`) stores the filtering and smoothing estimates, as well as the constructed model and input data. By default, `tempssm()` fits a first-order autoregressive model.

```
# model with first-order autoregressive component
res_ar1 <- tempssm(yamaguchi_sst) # (ar_order=1: default)
summary(res_ar1)
```

```
## tempssm summary
## -----
## Call:
## tempssm(temp_data = yamaguchi_sst)
##
## Model fit:
##   Likelihood type: marginal
##   Log-likelihood : -437.2
##   k               : 5
##   AIC             : 884.4
##   Converged       : TRUE
##
## Variance parameters:
##   Observation (H): 4.542027e-14
##   State (Q trend): 9.936473e-08
##   State (Q season): 8.408352e-56
##   State (Q ar): 0.3145626
##
## Components of auto-regression:
##   Order of AR: 1
##   Coefficient of AR1: 0.6202321
```

From the summary output, confirm that the model has converged (**Converged: TRUE**). The output also reports statistics such as the number of parameters (**k**), the log-likelihood, and the Akaike Information Criterion (**AIC**). The parameter estimates include the observation error variance (**H**), the process error variance of the long-term trend component (**Q trend**), the process error variance of the seasonal component (**Q season**), the process error variance of the autoregressive component, and the first-order autoregressive coefficient (**AR1**).

Examining the Autoregressive (AR) Order

Next, we examine how the autoregressive (AR) order affects model behavior. Specifically, we fit three models in which the order of the autoregressive component varies from 1 to 3, while all other model components, including the explicit seasonal cycle, are kept the same. Because the first-order model has already been fitted, we now fit the second- and third-order models.

```
# model with second-order autoregressive component
res_ar2 <- tempssm(yamaguchi_sst, ar_order=2)
summary(res_ar2)
```

```
## tempssm summary
## -----
## Call:
## tempssm(temp_data = yamaguchi_sst, ar_order = 2)
##
## Model fit:
```

```

## Likelihood type: marginal
## Log-likelihood : -434.26
## k : 6
## AIC : 880.51
## Converged : TRUE
##
## Variance parameters:
## Observation (H): 0.120908
## State (Q trend): 2.206262e-07
## State (Q season): 5.525557e-16
## State (Q ar): 0.1044204
##
## Components of auto-regression:
## Order of AR: 2
## Coefficient of AR1: 1.185818
## Coefficient of AR2: -0.4790522
# model with third-order autoregressive component
res_ar3 <- tempssm(yamaguchi_sst, ar_order=3)
summary(res_ar3)

## tempssm summary
## -----
## Call:
## tempssm(temp_data = yamaguchi_sst, ar_order = 3)
##
## Model fit:
## Likelihood type: marginal
## Log-likelihood : -434.09
## k : 7
## AIC : 882.19
## Converged : TRUE
##
## Variance parameters:
## Observation (H): 0.1123192
## State (Q trend): 2.242518e-07
## State (Q season): 9.553973e-10
## State (Q ar): 0.1245964
##
## Components of auto-regression:
## Order of AR: 3
## Coefficient of AR1: 1.07459
## Coefficient of AR2: -0.3334548
## Coefficient of AR3: -0.06336897

```

By comparing models with different AR orders, we examine how the representation of temporal dependence differs among the candidate models. Here, AIC is used as a relative measure for comparing these models.

Model Comparison Based on the Akaike Information Criterion (AIC)

The AIC of a fitted model can be obtained by applying `AIC()` to a `tempssm` object. By default, `tempssm()` fits the model with `marginal = TRUE`, using the marginal likelihood implemented by KFAS. This setting is stored in the fitted object and is carried forward to `logLik()` and `AIC()`. The diffuse likelihood remains available by explicitly setting `marginal = FALSE`.

For comparisons among several models, `compare_tempssm_aic()` provides a convenient interface. Before

returning a table containing AIC values, differences in AIC, Akaike weights, and related model information, the function checks the response series, observation period, convergence status, and likelihood type. All models being compared must use the same likelihood type so that their AIC values are calculated on a common basis.

```
# Extract AIC
```

```
AIC(res_ar1)
```

```
## [1] 884.3978
```

```
AIC(res_ar2)
```

```
## [1] 880.5134
```

```
AIC(res_ar3)
```

```
## [1] 882.1889
```

```
AIC_table_res <- compare_tempssm_aic(
  list(
    AR1 = res_ar1,
    AR2 = res_ar2,
    AR3 = res_ar3
  )
)
```

```
knitr::kable(AIC_table_res)
```

model	logLik	df	nobs	ob- served_n	start	end	fre- quency	likeli- hood	AIC	delta_AIC	weight
AR2	-	6	532	532	1982- 1	2026- 4	12	marginal	880.5134	0.000000	0.6344866
AR3	-	7	532	532	1982- 1	2026- 4	12	marginal	882.1889	1.675466	0.2745362
AR1	-	5	532	532	1982- 1	2026- 4	12	marginal	884.3978	3.884414	0.0909772


Among these candidate models, the AR2 model has the smallest AIC and therefore receives the strongest relative support according to this criterion. The following analyses use the AR2 model as the working model.

A smaller AIC does not, however, establish absolute model adequacy or guarantee better predictive performance. A more comprehensive assessment should also consider residual diagnostics and time-series cross-validation.

Plotting Long-Term Trend, Drift, Seasonal, and Autoregressive Components We extract the corresponding latent components from the state-space model and visualize the estimated long-term evolution of temperature levels and their rate of change (drift). Plotting seasonal variation and autoregressive dependence as well makes the underlying trend structure easier to examine.

```
# plot all components at once
```

```
plot(res_ar2)
```



The long-term trend in the upper-left panel indicates an increasing SST pattern over the study period. In this example, the average annual rate of SST increase over the study period is approximately 0.036 °C.

However, the rate of SST change itself also appears to vary over time. Focusing on the rate of change (drift) in the upper-right panel, the rate of change declines to zero or near zero during the 2000s and then increases after 2010. The shaded gray areas represent 95% confidence intervals for the estimated latent states and illustrate the uncertainty associated with each estimated component.

The standard plotting interface is `plot(res)`. The ggplot2-style interface `autoplot(res)` is also available and produces the same component plot by default. To obtain a ggplot object for an individual component, specify the `component` argument in `autoplot()` as follows.

```
# plot each of components at once
plt_level <- autoplot(res_ar2, component = c("level"))
plt_drift <- autoplot(res_ar2, component = c("drift"))
plt_season <- autoplot(res_ar2, component = c("season"))
plt_ar1 <- autoplot(res_ar2, component = c("ar1"))
```


Model Diagnostics

```
diag <- diagnose_residuals(res_ar2)
print(diag)
```

```
## # A tibble: 1 x 4
##   lb_stat lb_df lb_pvalue kurtosis
##   <dbl> <dbl>   <dbl>   <dbl>
## 1     3.68     2     0.158     3.66
```



```
plot_tempsm_residual_diagnostics(res_ar2)
```



In the model diagnostic plot, the upper panel shows the residual time series, the lower-left panel shows the residual autocorrelation plot (ACF plot), and the lower-right panel shows the residual frequency distribution. These plots should be checked for any notable residual patterns. The Ljung-Box test indicates no significant residual autocorrelation ($P = 0.996$).

Estimated Parameters and Latent-State Components

The long-term trend component and its rate of change (drift) can be extracted as `ts` objects as follows.

```
# Smoothing estimates
```

```
alpha_hat <- res_ar2$kfs$alphahat
head(alpha_hat)
```

```
##           level      slope sea_dummy1 sea_dummy2 sea_dummy3 sea_dummy4
## Jan 1982 18.95931 0.002015924 -4.388281 -1.989806 0.8433008 3.3260567
## Feb 1982 18.96132 0.002016062 -5.783414 -4.388281 -1.9898060 0.8433008
## Mar 1982 18.96334 0.002016254 -5.905548 -5.783414 -4.3882811 -1.9898060
## Apr 1982 18.96536 0.002016894 -4.594205 -5.905548 -5.7834141 -4.3882811
## May 1982 18.96737 0.002017326 -1.902724 -4.594205 -5.9055484 -5.7834141
## Jun 1982 18.96939 0.002017769 1.458513 -1.902724 -4.5942046 -5.9055484
##           sea_dummy5 sea_dummy6 sea_dummy7 sea_dummy8 sea_dummy9 sea_dummy10
## Jan 1982 6.1502216 7.6638746 5.1220111 1.4585130 -1.902724 -4.594205
## Feb 1982 3.3260567 6.1502216 7.6638746 5.1220111 1.458513 -1.902724
## Mar 1982 0.8433008 3.3260567 6.1502216 7.6638746 5.122011 1.458513
## Apr 1982 -1.9898060 0.8433008 3.3260567 6.1502216 7.663875 5.122011
## May 1982 -4.3882811 -1.9898060 0.8433008 3.3260567 6.150222 7.663875
```

```
## Jun 1982 -5.7834141 -4.3882811 -1.9898060 0.8433008 3.326057 6.150222
##          sea_dummy11      arima1      arima2
## Jan 1982 -5.905548 0.2080998 -0.06449083
## Feb 1982 -4.594205 0.2341012 -0.09969067
## Mar 1982 -1.902724 0.2821581 -0.11214667
## Apr 1982 1.458513 0.2875601 -0.13516843
## May 1982 5.122011 0.5397162 -0.13775631
## Jun 1982 7.663875 0.5449238 -0.25855222
```

```
# Smoothing estimate of level component
level_ts <- get_level_ts(res_ar2)

# Smoothing estimate of drift component
drift_ts <- get_drift_ts(res_ar2)

# Average drift rate per year across the full period
mean_drift_year <- mean(drift_ts)
print(mean_drift_year)
```

```
## [1] 0.03663453
```

The average annual rate of SST increase was estimated to be 0.0366 °C.

Exercise II: Applying a State-Space Model to a Temperature Time Series with an Exogenous Variable

Objective

We extend the state-space modeling framework to examine the effect of an exogenous factor on temperature variation. Specifically, we examine the effect of the Pacific Decadal Oscillation (PDO) as an exogenous variable on SST observed off Yamaguchi Prefecture, Japan.

Loading the PDO Index Dataset: PDO Index as an Exogenous Variable

- **Data:** Monthly Pacific Decadal Oscillation (PDO) index (JMA)
- **Period:** January 1901 to December 2025

The Pacific Decadal Oscillation (PDO) index is defined as the projection of monthly mean sea surface temperature (SST) anomalies onto the leading empirical orthogonal function (EOF) of SST variability over the North Pacific north of 20°N. The EOF is computed using SST anomalies for 1901–2000, defined relative to the 1901–2000 monthly climatology. To remove the global warming signal, the global-mean SST anomaly is subtracted from each grid point prior to the EOF analysis. In this package, we use the PDO index provided by the Japan Meteorological Agency (JMA), available at https://www.data.jma.go.jp/kaiyou/data/shindan/b_1/pdo/pdo.txt.

```
data(pdo) # load a ts object of NAO index
head(pdo)
```

```
##          Jan      Feb      Mar      Apr      May      Jun
## 1901  1.0040  0.7403  0.9011 -0.0109 -0.2325 -0.6810
```

Intersecting the Temperature and PDO Time Series

For state-space modeling with exogenous variables, all input time series must share a common and aligned time index. In this step, the temperature and PDO time series are restricted to their overlapping period by trimming the leading and trailing portions so that both datasets cover the same time span.

The function `tempssm::trim_ts_overlap()` is used to align the two `ts` objects on a shared timeline and returns time series containing only the common period.

```
# Generate an object on a shared timeline
yamaguchi_sst_trim <- trim_ts_overlap(yamaguchi_sst,
                                     pdo,
                                     temp_name = "Temp",
                                     exo_name="PDO")$temperature

pdo_trim <- trim_ts_overlap(yamaguchi_sst,
                           pdo,
                           temp_name = "Temp",
                           exo_name="PDO")$exogenous

start(yamaguchi_sst_trim)

## [1] 1982    1
end(yamaguchi_sst_trim)

## [1] 2025   12
start(pdo_trim)

## [1] 1982    1
end(pdo_trim)

## [1] 2025   12
```


Plotting Time Series of SST and the PDO Index

We visualize the time-series structures of SST and the PDO index.

```
plt_yamaguchi_sst_trim <- forecast::autoplot(yamaguchi_sst_trim) +
  labs(y = expression(Temp.~(degree*C)),
       x = "Time (year)") +
  ggtitle("Monthly SST off Yamaguchi, Japan") +
  theme_classic()

plt_pdo <- forecast::autoplot(pdo_trim) +
  labs(x = "Time (year)", y = "PDO index") +
  ggtitle("PDO index") +
  theme_classic()

plt_yamaguchi_sst_trim + plt_pdo + patchwork::plot_layout(ncol=1)
```



The PDO index data used in this exercise contain no missing values. When using your own exogenous-variable data, however, make sure that the exogenous variables do not contain missing values before fitting the model. `tempssm()` does not allow missing values in exogenous variables, and model fitting will stop with an error if such data are supplied. If necessary, apply appropriate preprocessing, such as imputation, before the analysis.

In contrast, missing values in the dependent temperature time series are allowed. In the state-space model, these missing values are treated as unobserved responses, and model estimation proceeds through Kalman filtering and smoothing.

Applying a Model Without an Exogenous Variable

We first apply a baseline state-space model that does not include any exogenous variables. This model serves as a reference case in which temperature variation is explained solely by the latent long-term trend, seasonal cycle, and autoregressive dependence. Although the same model was already fitted in Exercise I, here we explicitly specify the dependent-variable argument and store the returned object as `res_without`.

```
res_without <- tempssm(temp_data = yamaguchi_sst_trim, ar_order = 2)
summary(res_without)
```

```
## tempssm summary
## -----
## Call:
## tempssm(temp_data = yamaguchi_sst_trim, ar_order = 2)
##
## Model fit:
##   Likelihood type: marginal
##   Log-likelihood : -432.73
##   k               : 6
##   AIC             : 877.47
```

```
## Converged      : TRUE
##
## Variance parameters:
## Observation (H): 0.1207584
## State (Q trend): 2.163319e-07
## State (Q season): 2.411001e-20
## State (Q ar): 0.1071204
##
## Components of auto-regression:
## Order of AR: 2
## Coefficient of AR1: 1.17824
## Coefficient of AR2: -0.4720514
```

The summary output should be the same as `summary(res_ar2)` from Exercise I. Prior examination of autoregressive orders from AR(1) to AR(3) indicated that AR(2) provided the best model fit not only for the baseline model but also for the model including the exogenous PDO index. The details are omitted from this manual, but interested users are encouraged to examine alternative AR orders in their own analyses.

This baseline model provides a useful benchmark for evaluating the additional explanatory power of the PDO index introduced in the next section.

Applying a Model With an Exogenous Variable

We construct a model that uses the PDO index as a univariate exogenous variable. In the `tempssm()` function, the PDO data (`pdo_trim`) are supplied to the exogenous-variable argument (`exo_data`). The temperature time series (`yamaguchi_sst_trim`) is supplied to the dependent-variable argument (`temp_data`), which was implicit in the previous examples.

```
res_with <- tempssm(temp_data = yamaguchi_sst_trim,
                    exo_data = pdo_trim,
                    ar_order = 2)
summary(res_with)

## tempssm summary
## -----
## Call:
## tempssm(temp_data = yamaguchi_sst_trim, exo_data = pdo_trim,
##       ar_order = 2)
##
## Model fit:
## Likelihood type: marginal
## Log-likelihood : -390.91
## k              : 7
## AIC            : 795.83
## Converged      : TRUE
##
## Variance parameters:
## Observation (H): 0.05749006
## State (Q trend): 3.832674e-19
## State (Q season): 1.5571e-54
## State (Q ar): 0.1792551
##
## Components of auto-regression:
## Order of AR: 2
## Coefficient of AR1: 0.8656257
## Coefficient of AR2: -0.1615369
```

```
## Exogenous variable    PDO
## Estimated coefficient   -0.4406609
## Lower CI   -0.5296035
## Upper CI   -0.3517183
```

We then examine the model estimates. The estimated coefficient for the exogenous PDO index is negative (-0.44), and its 95% confidence interval does not include zero

-0.53, -0.35

. This indicates a statistically significant negative relationship between local SST off Yamaguchi Prefecture and PDO variability.

Specifically, after accounting for the underlying long-term trend, seasonal cycle, and autoregressive dependence, the model suggests that a one-unit increase in the PDO index is associated with an average decrease of approximately 0.44 °C in monthly SST. This result implies that positive phases of the PDO systematically contribute to cooler conditions at the study site.

Importantly, the effect of this exogenous variable is identified as an additional factor beyond the internal dynamics of the temperature time series, rather than as an artifact of improved representation of the long-term trend or temporal dependence. Compared with the baseline model without exogenous variables, the statistical significance of the PDO coefficient indicates that the PDO acts as an independent large-scale climate driver influencing local temperature variation.

Next, we compare the overall goodness of fit between the two models using the Akaike Information Criterion (AIC).

Model Comparison Based on AIC

Model selection criteria such as AIC provide quantitative measures of model adequacy that balance goodness of fit against model complexity. Lower AIC values indicate a more parsimonious model with stronger support from the data.

```
AIC_without <- AIC(res_without)
AIC_with <- AIC(res_with)

models_AICs <- tibble(
  model = c("Without", "With"),
  AIC = c(AIC_without, AIC_with),
  delta_AIC = min(AIC) - AIC
)

models_AICs %>% knitr::kable()
```

model	AIC	delta_AIC
Without	877.4658	-81.63906
With	795.8268	0.00000

The model including the PDO index as an exogenous variable exhibits a substantially lower AIC than the baseline model without exogenous variables, indicating a clear improvement in overall fit. This result supports the conclusion that explicitly accounting for PDO variability improves the statistical description of the temperature time series beyond what is achieved by internal dynamics alone.

It should be noted that AIC-based comparisons are meaningful only among models fitted to the same time series over an identical period. Moreover, in state-space models, AIC differences reflect both regression effects and changes in the stochastic structure of latent components. Therefore, AIC should be used as a

complementary diagnostic alongside parameter estimates and their uncertainty, rather than as the sole basis for inference.

To further evaluate model robustness and predictive performance, we use time-series cross-validation as described below.

Time-Series Cross-Validation (tsCV)

Time-series cross-validation (tsCV) evaluates model performance based on out-of-sample prediction errors while respecting the temporal ordering of the data. It therefore provides an additional and independent perspective on model adequacy.

In time-series cross-validation, the model is repeatedly fitted to expanding or sliding training windows, and predictive performance is evaluated on subsequent observations. Unlike random cross-validation, this procedure prevents information leakage from the future to the past and is therefore suitable for model validation with temporal data.

Here, by comparing cross-validation metrics for models with and without the exogenous PDO variable, we assess whether the improvement suggested by AIC is also reflected in out-of-sample predictive performance.

```
# (Optional) Load packages for parallel processing.
# These are only required if you want to enable parallel execution:
# - future: controls how parallel workers are launched
# - future.apply: provides parallelized apply functions
# - progressr: optional progress bar support
library(future.apply)
library(future)
library(progressr)

if (interactive()) {
  future::plan(multisession)
} else {
  future::plan(sequential)
}

# ts cross-validation of the model without exogenous variables

## Generate a list of training and test datasets with their indices
# Procedure for constructing year-based time-series cross-validation folds:
#
# First training data: January 1982-December 2017;
# First test data: one year starting from January 2018
#
# Second training data: January 1983-December 2018;
# Second test data: one year starting from January 2019
# ...
#
# Eighth training data: January 1989-December 2024;
# Eighth test data: one year starting from January 2025
#
# These folds are automatically generated by the ts_train_test_split() function.

# Generate training and test dataset for model without PDO index
folds_without <- ts_train_test_split(
  temp_data = yamaguchi_sst_trim,
```

```

exo_data = NULL,
initial = 432, # 432 monthly observations from Jan 1982 to Dec 2017
horizon = 12, # forecast 12 monthly time-series
step = 12, # training data is moved by 12 months (1 years) steps
fixed_window = TRUE,
allow_partial = FALSE
)

# Generate training and test dataset for model with PDO index
folds_with <- ts_train_test_split(
  temp_data = yamaguchi_sst_trim,
  exo_data = pdo_trim,
  initial = 432, # 432 monthly observations from Jan 1982 to Dec 2017
  horizon = 12, # forecast 12 monthly time-series
  step = 12, # training data is moved by 12 months (10 years) steps
  fixed_window = TRUE,
  allow_partial = FALSE
)

# *****
# Executing time-series cross validation

#-----
# Model without the exogenous variable of PDA index

# Single processing
# cv_without_results <- ts_cv_run(folds_without, ar_order = 2, use_season = TRUE)

# Parallel processing
cv_without_results <- future.apply::future_lapply(
  folds_without,
  ts_cv_run_fold,
  ar_order = 2,
  use_season = TRUE,
  future.seed = TRUE
)

# Computing assessment indexes
metrics_without <- lapply(cv_without_results, compute_cv_metrics)

# tidy summary
cv_without_tbl <- ts_cv_collect(cv_without_results, metrics_without) %>%
  mutate(Model="Without")

#-----
# Model with the exogenous variable of PDA index

# Single processing
#cv_with_results <- ts_cv_run(folds_with, ar_order = 2, use_season = TRUE)

# Parallel processing

```



```

cv_with_results <- future.apply::future_lapply(
  folds_with,
  ts_cv_run_fold,
  ar_order = 2,
  use_season = TRUE,
  future.seed = TRUE
)

# Computing assessment indexes
metrics_with <- lapply(cv_with_results, compute_cv_metrics)

# tidy summary
cv_with_tbl <- ts_cv_collect(cv_with_results, metrics_with) %>%
  mutate(Model="With")

cv_tbl <- bind_rows(cv_without_tbl, cv_with_tbl)


plt_MAE <- ggplot(data=cv_tbl,
  aes(x=Model,y=MAE)) +
  geom_boxplot()

plt_MASE_naive <- ggplot(data=cv_tbl,
  aes(x=Model,y=MASE_naive)) +
  geom_boxplot()

plt_MASE_seasonal <- ggplot(data=cv_tbl,
  aes(x=Model,y=MASE_seasonal)) +
  geom_boxplot()

plt_tsCV <- plt_MAE + plt_MASE_naive + plt_MASE_seasonal + patchwork::plot_layout(nrow=1)
plot(plt_tsCV)

```



The tsCV results show that the model with the exogenous variable performs better than the model without it. In this tutorial, the number of tsCV iterations is set to eight to reduce execution time. In actual validation, please use a sufficient number of iterations.

Taken together, the results consistently support including the PDO index as an exogenous variable in the state-space model. The PDO coefficient is statistically significant, indicating a clear effect on temperature variation. The improvement in AIC also indicates better overall model fit. Furthermore, time-series cross-validation confirms that this improvement translates into better out-of-sample predictive performance.

These complementary lines of evidence provide robust and multifaceted support for adopting the exogenous-variable model.

Comparing Estimation Patterns Between Models With and Without the Exogenous Variable

We compare how the estimated patterns of the long-term trend and its rate of change differ between the models with and without the exogenous variable.



```
plt_level_without_ts <- autoplot(res_without,
                                component=c("level")
                                ) +
  labs(title="Model without the PDO index") +
  theme_classic()

plt_drift_without_ts <- autoplot(res_without,
                                component=c("drift")
                                ) +
  labs(title="Model without the PDO index") +
  theme_classic()

plt_level_drift_without_ts <- plt_level_without_ts +
```

```
plt_drift_without_ts +
  patchwork::plot_layout(ncol=1)

plot(plt_level_drift_without_ts)
```





```
plt_level_with_ts <- autoplot(res_with,
                             component=c("level")
                             ) +
  labs(title="Model with the PDO index") +
  theme_classic()

plt_drift_with_ts <- autoplot(res_with,
                             component=c("drift")
                             ) +
  labs(title="Model with the PDO index") +
  theme_classic()

plt_level_drift_with_ts <- plt_level_with_ts +
  plt_drift_with_ts +
  patchwork::plot_layout(ncol=1)

plot(plt_level_drift_with_ts)
```



```
# Smoothing estimate of drift component
mean_drift_year_without <- get_drift_ts(res_without) %>%
  mean()
print(mean_drift_year_without)

## [1] 0.03601619

mean_drift_year_with <- get_drift_ts(res_with) %>%
  mean()
print(mean_drift_year_with)

## [1] 0.01142969
```

The gray areas in the plots above show 95% confidence intervals.

In the model without an exogenous variable, the plots of the long-term trend and its rate of change show a wavelike pattern. In contrast, in the model with the exogenous variable, the long-term trend and its rate of change show a more stable pattern. This suggests that the wavelike pattern reflects the influence of the exogenous PDO variable.

The estimated annual rate of SST change over the observation period is 0.0360 in the model without the exogenous variable and 0.0114 in the model with the exogenous variable.

In this case, the model incorporating the PDO index is considered preferable for accurately estimating the SST trend. Thus, the SST warming rate, which was estimated somewhat higher in the model without the exogenous variable, is estimated more conservatively after introducing the model with the exogenous variable.

Using Your Own Data

Quick Example

`my_ts` is a user-supplied `ts` object containing temperature time-series data. It can be passed directly to `tempssm()`.

```
## example of ts object (not run)
## my_ts <- ts(rnorm(100), frequency = 12)
res <- tempssm(my_ts)
```

Preparing Your Own Dataset

CSV Format

Prepare a CSV file of monthly temperature time series with the following columns: - Year - Month - Temp

Example:

```
Year,Month,Temp
2010,8,13.6
2010,9,6.8
2010,10,NA
2010,11,-1.4
...
```

- Use NA for missing temperature values, and always keep the corresponding Year and Month entries.
- The CSV file must be comma-separated and UTF-8 encoded.

Converting CSV to a Time Series

An example CSV file included in this package is available in `inst/extdata`. The example dataset contains monthly air temperature observations from Mt. Akadake, Hokkaido, Japan (1,840 m). The original data are available from the Monitoring Sites 1000 Project of the Ministry of the Environment of Japan (KOZ01.zip, <https://www.biodic.go.jp/moni1000/findings/data/index.html>).

```
path <- system.file("extdata", "example_monthly_temp.csv", package = "tempssm")
akadake_temp <- tempssm::read_monthly_temp_ts(path)
head(akadake_temp)
```

```
##      Jan Feb Mar Apr May Jun Jul   Aug   Sep   Oct   Nov   Dec
## 2010              13.6   6.8   0.2  -6.8 -12.5
## 2011 -18.8
```

The `read_monthly_temp_ts()` function internally uses the base R `ts()` function to create a time-series object. If finer control is needed, you can construct the `ts` object manually.

References

The statistical modeling framework implemented in `tempssm` is based on the methodology described in Baba et al. (2024), and the accompanying supplementary materials and code repository served as the initial source for development of this package.

Baba, S., Ishii, H., and Yoshiyama, T. (2024). Estimating sea temperature trends using a linear Gaussian state-space model in Jogashima, Kanagawa, Japan. *Bulletin of the Japanese Society of Fisheries Oceanography*, 88(3), 190–199. (In Japanese with an English abstract.) https://doi.org/10.34423/jsfo.88.3_190

Baba, S. (2024). Supplementary code and test data for estimating sea temperature trends using a linear Gaussian state-space model. GitHub repository: <https://github.com/logics-of-blue/sea-temperature-trend-jogashima>

Helske, J. (2017). KFAS: Exponential Family State Space Models in R. *Journal of Statistical Software*, 78(10), 1–39. <https://doi.org/10.18637/jss.v078.i10>

Appendix: Utility Functions

The following utility functions are provided to support data preparation and exploratory analysis.

1. `read_monthly_temp_ts()`

The `read_monthly_temp_ts()` function converts monthly temperature data stored in a CSV file into an R `ts` object. By enforcing a simple and consistent data format, it is intended to make externally prepared time-series data easier to import into `tempssm`.

Example

For monthly temperature data, prepare a CSV file with a header row. By default, the column names must be `Year`, `Month`, and `Temp`, where `Temp` represents the observed temperature value.

```
Year,Month,Temp
2001,1,10.4
2001,2,8.2
2001,3,NA
2001,4,13.6
2001,5,16.1
...
```

- Use NA for missing temperature values, and always keep the corresponding Year and Month entries.
- The CSV file must be comma-separated and UTF-8 encoded.

The following example uses a sample CSV file included in the package.

```
tmp_csv <- tempfile(fileext = ".csv")
writeLines(
  c("Year,Month,Temp",
    "2001,1,10.4",
    "2001,2,8.2",
    "2001,3,NA",
    "2001,4,13.6",
    "2001,5,16.1"),
  tmp_csv
)

# Read the CSV file and convert to a monthly ts object
temp_ts <- read_monthly_temp_ts(tmp_csv)
```

2. `convert_monthly_df_to_ts()`

The `convert_monthly_df_to_ts()` function converts a data frame containing monthly temperature data into an R `ts` object. It is designed to support workflows in which temperature data are first imported or prepared as a data frame and then used for time-series analysis.

The input data frame is expected to contain at least two columns: a date column (`Date`) and a temperature column (`Temp`). For the date column (`Date`), assign an arbitrary day of the month, such as the first day, so

that the dates correspond to monthly aggregated temperature values.

Example

```
# Create a data frame of monthly temperature data
df <- data.frame(
  Date = as.Date(c(
    "2001-01-01",
    "2001-02-01",
    "2001-03-01",
    "2001-04-01",
    "2001-05-01")
  ),
  Temp = c(10.4, 8.2, NA, 13.6, 16.1)
)

# Convert to a monthly ts object
temp_ts <- convert_monthly_df_to_ts(df)
head(temp_ts)
```

```
##           Jan  Feb  Mar  Apr  May
## 2001 10.4   8.2   NA 13.6 16.1
```

3. get_jma_sst_ts()

The `get_jma_sst_ts()` function downloads daily mean sea surface temperature (SST) data for Japanese coastal waters provided by the Japan Meteorological Agency (JMA). It aggregates the daily values into monthly means and returns the resulting time series as an object of class `ts`.

Example

The following example demonstrates how to download SST data for the southern coastal waters of Ibaraki Prefecture, Japan.

The argument `sea_area_id` specifies the numeric identifier of a sea area defined by JMA. For example, 138 corresponds to the coastal waters off southern Ibaraki. A list of available sea area IDs and their corresponding regions is provided by JMA at:

https://www.data.jma.go.jp/kaiyou/data/db/kaikyo/series/engan/eg_areano.html

```
sst_138_ts <- get_jma_sst_ts(sea_area_id = 138)
head(sst_138_ts)
```

```
##           Jan      Feb      Mar      Apr      May      Jun
## 1982 15.04419 14.22500 13.63903 15.31933 17.52258 19.52300
```

4. compute_temp_anomaly()


The `compute_temp_anomaly()` function computes anomalies from a time series provided as an R `ts` object. Anomalies are calculated by subtracting the corresponding long-term mean for each periodic unit, such as month, from each observation. This transformation is useful for examining interannual variability after removing periodic effects.

The reference period used to compute anomalies can be specified via a function argument. By default, the climatology is computed over the entire available time series (see `?compute_temp_anomaly`).

This transformation is useful for exploratory analysis and modeling applications that focus on departures from typical seasonal conditions.

Example

```
# Generate temperature anomalies
data(niigata_sst)
niigata_sst_anomaly <- compute_temp_anomaly(niigata_sst)
plt_niigata_sst_anomaly <- forecast::autoplot(niigata_sst_anomaly) +
  labs(y = expression(Temp.~(degree*C)),
       x = "Time (year)")
plot(plt_niigata_sst_anomaly)
```



5. compute_monthly_climatology()

The `compute_monthly_climatology()` function computes the climatological mean seasonal cycle from a `ts` object by averaging each seasonal value across years. It is primarily intended for exploratory analysis and visualization of seasonal structure in temperature time series.

Example

```
data(niigata_sst)
monthly_seasonal_cycle_niigata_sst <- compute_monthly_climatology(niigata_sst)
summary(monthly_seasonal_cycle_niigata_sst)
```

```
##      Month      Temperature
## Min.   : 1.00   Min.    : 9.365
## 1st Qu.: 3.75   1st Qu.:11.169
## Median : 6.50   Median :16.318
## Mean   : 6.50   Mean    :17.036
## 3rd Qu.: 9.25   3rd Qu.:22.294
## Max.   :12.00   Max.    :26.873
```



```
plt_monthly_seasonal_cycle_niigata_sst <- ggplot(
  data = monthly_seasonal_cycle_niigata_sst,
  aes(x = Month, y = Temperature)
) +
  geom_point(size = 2) +
  geom_line(linetype = "dashed") +
  labs(
    title = "Monthly seasonal cycle of temperature",
    x = "Month",
    y = expression(Temp.~(degree*C))
  ) +
  scale_x_continuous(
    breaks = 1:12,
    labels = sprintf("%02d", 1:12)
  ) +
  theme_classic()

plot(plt_monthly_seasonal_cycle_niigata_sst)
```

