

tempssm: State-space modeling for temperature time series in R

Version 0.3.0

Akira S. Hirao, Shinya Baba, and Momoko Ichinokawa

2026-07-23

Summary

R package **tempssm** provides tools for state-space analysis of temperature time series, focusing on linear Gaussian state-space models estimated via Kalman filtering and smoothing as implemented in the KFAS package (Helske, 2017).

Key features

- Designed for temperature time series with arbitrary seasonal frequencies; currently validated primarily on monthly data
- Estimates latent states using linear Gaussian state-space models combined with Kalman filtering and smoothing
- Models temperature dynamics as a sum of interpretable latent components, including a long-term trend, seasonal cycle, autoregressive structure, and optional exogenous effects
- Allows users to specify an arbitrary order of the autoregressive component (default: AR(1))
- Implements time-series cross-validation for model evaluation

Prior Art and Scope

tempssm provides a domain-focused workflow for analyzing temperature time series with linear Gaussian state-space models. It brings together model construction, component extraction, uncertainty summaries, residual diagnostics, visualization, and time-series cross-validation in a single R package interface tailored to temperature applications.

The package builds on established statistical methodology, including linear Gaussian state-space modeling, Kalman filtering, and Kalman smoothing. Model estimation is handled through the KFAS package, which provides a general framework for state-space models in R.

The initial implementation was adapted from the supplementary code provided by Baba (2024), accompanying Baba et al. (2024), which analyzed sea temperature trends using a linear Gaussian state-space model. The supplementary code is publicly available at:

<https://github.com/logics-of-blue/sea-temperature-trend-jogashima>

Compared with that prior implementation, **tempssm** extends the workflow into a reusable R package interface with input validation, documented S3 methods, tests, diagnostics, cross-validation utilities, and examples for broader temperature time-series analysis.

The next section describes the state-space model used by **tempssm** and clarifies how the trend, seasonal, autoregressive, and exogenous components are represented.

State-Space Model in tempssm

We consider an extended version of the Basic Structural Time Series Model (BSTSM) to describe temperature time series with explicit long-term trend, seasonal variability, and auto-regressive component. The observation equation is given by

$$y_t = \alpha_t + v_t. \quad (1)$$

where t denotes the time index, y_t is the observed temperature, α_t is the latent state, and v_t is the observation error.

The latent state is decomposed as

$$\alpha_t = \mu_t + s_t + r_t + E_t. \quad (2)$$

where μ_t is the level (trend), s_t the seasonal component, r_t a stationary autoregressive component, and E_t the contribution of exogenous variables.

The level component follows a second-order stochastic process:

$$\mu_t = 2\mu_{t-1} - \mu_{t-2} + \zeta_t, \quad \zeta_t \sim \mathcal{N}(0, \sigma_\zeta^2). \quad (3)$$

where ζ_t is the process error of the level component.

The seasonal component is modeled with a sum-to-zero constraint:

$$s_t = - \sum_{i=t-f}^{t-1} s_i + \omega_t, \quad \omega_t \sim \mathcal{N}(0, \sigma_\omega^2). \quad (4)$$

where ω_t is the process error of the seasonal component and f denotes the seasonal frequency (e.g., $f = 12$ for monthly data). This formulation ensures that seasonal effects are identifiable and comparable across different temporal resolutions. By enforcing the sum-to-zero constraint over one full seasonal cycle, the seasonal component captures recurring deviations from the underlying trend without introducing long-term drift. In models without a seasonal component, the seasonal term s_t is set to zero and omitted from the state equation.

The autoregressive component follows an l th-order autoregressive (AR) process:

$$r_t = \phi_1 r_{t-1} + \phi_2 r_{t-2} + \dots + \phi_l r_{t-l} + \tau_t, \quad \tau_t \sim \mathcal{N}(0, \sigma_\tau^2). \quad (5)$$

Here, τ_t denotes the process error of the autoregressive component. The process error terms (ζ_t , ω_t , and τ_t) are assumed to be mutually independent and normally distributed.

The exogenous component is defined as

$$E_t = \beta_1 x_{1,t} + \beta_2 x_{2,t} + \dots + \beta_m x_{m,t}. \quad (6)$$

The exogenous term E_t represents the influence of external factors, which may include large-scale climate indices, regional environmental variables, or other physically motivated predictors relevant to the observed temperature time series. Each exogenous variable enters the model linearly through a time-invariant regression coefficient ($\beta_1, \beta_2, \dots, \beta_m$), allowing both the magnitude and direction of its contribution to be estimated jointly with the latent state components.

When applying the model to observational data, the total number of parameters to be estimated (k) depends on the order of autoregressive component and the number of exogenous variables included. For example, in a model with a second-order autoregressive component and no exogenous covariates, six parameters are estimated: the observation error variance, the process error variances associated with the level, seasonal, and autoregressive components, and the first- and second-order autoregressive coefficients (ϕ_1 and ϕ_2). In the general case with an l -order autoregressive component and m exogenous variables, a total number of parameters is $k = 4 + l + m$.

The core implementation of parameter estimation procedure is based on the supplementary code provided by Baba et al. (2024). Parameter estimation is performed using a two-step optimization strategy recommended

by Helske (2017). In the first step, model parameters are estimated using the Nelder-Mead method (Nelder & Mead, 1965) with user-specified initial values. The resulting estimates are then used as initial values in a second optimization step based on the BFGS algorithm (Shanno, 1970). The main model-fitting function, `tempssm()`, returns both filtering and smoothing estimates. Unless otherwise stated, all results presented in this vignette are based on the smoothed estimates.

How to use

Set Environment

Load the following libraries for executing ‘How to use’.

```
## Set libraries
library(tempssm)
```

Input Data Format

Input data for `tempssm` must be supplied as an R `ts` object, which represents a regularly spaced time series (see `?stats::ts` or <https://stat.ethz.ch/R-manual/R-devel/library/stats/html/ts.html>).

To support data preparation, the package includes utility functions that convert external observational data into `ts` objects suitable for model fitting (see Appendix).

Practice: Applying State-Space Model to a Univariate Temperature Time Series

Objective

The objective of this practice is to demonstrate the basic application of a linear Gaussian state-space model to a univariate temperature time series. This example serves as an introduction to the modeling framework and highlights the role of autoregressive dynamics without the inclusion of exogenous variables.

Loading the Dataset

A sample sea surface temperature (SST) dataset is included in the package.

- **Dataset:** Monthly sea surface temperature (SST) off Niigata, Japan
- **Unit:** Degrees Celsius
- **Period:** February 2002 to December 2023

This dataset is derived from observations archived at Japan Oceanographic Data Center (JODC), Hydrographic and Oceanographic Department, Japan Coast Guard. Original daily SST data were obtained from <https://www.jodc.go.jp/jodcweb/JDOSS/index.html> and aggregated into monthly means.

```
data(niigata_sst) # load a ts object of SST off Niigata
head(niigata_sst)
```

```
##           Jan           Feb           Mar           Apr           May           Jun
## 2002  9.951613  8.332143  9.348387 11.713333 14.529032 18.906667
```

```
summary(niigata_sst)
```

```
##      Temp
## Min.   : 7.707
## 1st Qu.:11.217
## Median :16.345
## Mean   :17.033
```

```
## 3rd Qu.:22.787
## Max.    :28.897
## NAs     :2
```


The dataset includes two missing observations. Even if missing observations was in your dataset included, there are retained and handled explicitly within the state-space modeling framework.

Plotting the Monthly SST Time Series

We begin by visualizing the monthly SST time series to examine its overall structure, including apparent trends, seasonal variability, and missing observations.

```
plt_niigata_sst <- forecast::autoplot(niigata_sst) +
  ggplot2::labs(
    y = expression(Temperature ~ (degree * C)),
    x = "Time (year)"
  ) +
  ggplot2::ggtitle("Monthly SST off Niigata, Japan") +
  ggplot2::theme_classic()

plot(plt_niigata_sst)
```



The overall mean SST is approximately 17 °C, and a clear seasonal pattern is visible. The series contains two missing observations, in September 2005 and February 2006. Although SST appears to be higher near the end of the series than near the beginning, year-to-year variability is also evident, making the long-term trend difficult to assess from the raw time series alone.

Applying a Linear Gaussian State-Space Model

When a `ts` object containing temperature time-series data (here, `niigata_sst`) is passed to the core function `tempssm()`, model construction and parameter estimation are performed together. The returned S3 object of class `tempssm` (here, `res`) stores the filtering and smoothing estimates, as well as the constructed model and input data. By default, `tempssm()` fits a first-order autoregressive model.

```
# model with first-order autoregressive component
res <- tempssm(niigata_sst) # AR(1), the default model
summary(res)
```

```
## tempssm summary
## -----
## Call:
## tempssm(temp_data = niigata_sst)
##
## Model fit:
##   Likelihood type: marginal
##   Log-likelihood : -249.8
##   k               : 5
##   Diffuse states : 13
##   Converged       : TRUE
##
## Variance parameters:
##   Observation (H): 0.005985637
##   State (Q trend): 1.268117e-07
##   State (Q season): 0.001346138
##   State (Q ar): 0.4097883
##
## Components of auto-regression:
##   Order of AR: 1
##   Coefficient of AR1: 0.7442999
```

From the summary output, confirm that the model has converged (Converged: TRUE). The output also reports statistics such as the number of parameters (`k`), the log-likelihood, the likelihood type, and the number of diffuse initial states. The parameter estimates include the observation error variance (`H`), the process error variance of the long-term trend component (`Q trend`), the process error variance of the seasonal component (`Q season`), the process error variance of the autoregressive component, and the first-order autoregressive coefficient (`AR1`).

The log-likelihood and the associated number of estimated parameters can also be extracted directly from the fitted `tempssm` object using `logLik()`.

```
ll <- logLik(res)
ll

## 'log Lik.' -249.7962 (df=5)
attr(ll, "df") # number of parameters
```


```
## [1] 5
```

By default, `tempssm()` uses the KFAS marginal likelihood for parameter estimation. The selected likelihood type is retained for `logLik()` and `summary()`. The package intentionally does not compute AIC for `tempssm` objects. The log-likelihood and parameter count remain available through `logLik()` for users who need them for their own model-assessment workflows. The diffuse likelihood remains available by fitting the model with `marginal = FALSE`.

Plotting Level, Drift, Seasonal, and Auto-Regressive Components

We visualize the estimated long-term evolution of temperature levels and their rates of change (drift) by extracting the corresponding latent components from the state-space model. Additionally, seasonal variability and autoregressive dependence are plotted out, allowing the underlying trend behavior to be examined more clearly.

```
# plot all components at once  
plot(res)
```



The level component shows a persistent upward trend in sea surface temperature over the study period, while the drift component indicates a relatively stable positive rate of change. The shaded gray areas represent 95% confidence intervals for the estimated latent states, illustrating the uncertainty associated with each of estimated components.


Although the observed time series contains missing values, the state-space framework allows latent states to be estimated for the entire time span, including unobserved periods.

The standard plotting interface is `plot(res)`. The ggplot2-style interface `autoplot(res)` is also available and produces the same component plot by default. It returns a faceted `ggplot` object, so selected components can be stored and customized with standard ggplot2 layers, for example, `autoplot(res, component = c("level", "drift")) + ggplot2::theme_bw()`.

Simple Model Diagnostics

The package provides diagnostic tools for checking whether the fitted model has left notable structure in the residuals. In particular, residual time-series plots, residual autocorrelation, residual distributions, and Ljung-Box tests can be used to assess remaining temporal dependence and departures from the Gaussian error assumption.

```
plot_tempssm_residual_diagnostics(res)
```



In the model diagnostic plot, the upper panel shows the residual time series, the lower-left panel shows the residual autocorrelation plot (ACF plot), and the lower-right panel shows the residual frequency distribution. These plots should be checked for any notable residual patterns.

```
diag <- diagnose_residuals(res)
print(diag)
```

```
## # A tibble: 1 x 4
##   lb_stat lb_lag lb_pvalue kurtosis
##   <dbl> <dbl>   <dbl>   <dbl>
## 1    10.7    12     0.558     3.06
```

The `lb_stat`, `lb_lag`, and `lb_pvalue` columns correspond to the Ljung-Box test statistic, the lag used in the test, and its P-value, respectively. For monthly time series, `diagnose_residuals()` uses lag 12 by default. In this example, the Ljung-Box test indicated no significant residual autocorrelation up to lag 12 ($P > 0.05$).

Estimated Parameters and Latent-State Components

The long-term trend component and its rate of change (drift) can be extracted as ts objects as follows.

```
# Smoothing estimates
alpha_hat <- res$kfs$alphahat
head(alpha_hat)
```

```
##           level      slope sea_dummy1 sea_dummy2 sea_dummy3 sea_dummy4
## Jan 2002 16.3944 0.005600279 -6.624679 -3.337435  0.6067453  4.7593086
## Feb 2002 16.4000 0.005600420 -7.676822 -6.624679 -3.3374348  0.6067453
```

```
## Mar 2002 16.4056 0.005600414 -7.346316 -7.676822 -6.6246785 -3.3374348
## Apr 2002 16.4112 0.005600310 -5.476554 -7.346316 -7.6768220 -6.6246785
## May 2002 16.4168 0.005600334 -2.217000 -5.476554 -7.3463157 -7.6768220
## Jun 2002 16.4224 0.005600466 2.468021 -2.217000 -5.4765545 -7.3463157
##          sea_dummy5 sea_dummy6 sea_dummy7 sea_dummy8 sea_dummy9 sea_dummy10
## Jan 2002 8.5280151 9.9007960 6.4159192 2.4680212 -2.217000 -5.476554
## Feb 2002 4.7593086 8.5280151 9.9007960 6.4159192 2.468021 -2.217000
## Mar 2002 0.6067453 4.7593086 8.5280151 9.9007960 6.415919 2.468021
## Apr 2002 -3.3374348 0.6067453 4.7593086 8.5280151 9.900796 6.415919
## May 2002 -6.6246785 -3.3374348 0.6067453 4.7593086 8.528015 9.900796
## Jun 2002 -7.6768220 -6.6246785 -3.3374348 0.6067453 4.759309 8.528015
##          sea_dummy11      arima1
## Jan 2002 -7.346316 0.17523167
## Feb 2002 -5.476554 -0.37744125
## Mar 2002 -2.217000 0.28684018
## Apr 2002 2.468021 0.76796641
## May 2002 6.415919 0.33018674
## Jun 2002 9.900796 0.00904254
```

```
# Smoothing estimate of level component
level_ts <- get_level_ts(res)

# Smoothing estimate of drift component
drift_ts <- get_drift_ts(res)

# Average drift rate per year across the full period
mean_drift_year <- mean(drift_ts)
print(mean_drift_year)
```

```
## [1] 0.05259212
```

Average annual increase in SST is approximately 0.05 °C.

Short-Term Prediction

A fitted `tempssm` object can also be passed to `predict()` to obtain short-term predictions. By default, `predict(res)` returns a one-step-ahead prediction beyond the end of the observed series. This is useful for visual checks of how the fitted model extrapolates the estimated level, seasonal, and autoregressive components.

```
pred_1 <- predict(res)
pred_1
```

```
##          Jan
## 2024 10.61997
```

Predictions for multiple future time points can be requested by setting the `n.ahead` argument.

```
pred_12 <- predict(res, n.ahead = 12)
pred_12
```

```
##          Jan          Feb          Mar          Apr          May          Jun          Jul
## 2024 10.619966 9.500226 10.198500 11.757055 15.414159 19.704126 24.218777
##          Aug          Sep          Oct          Nov          Dec
## 2024 27.331048 25.941035 22.395225 18.262379 14.141777
```

These predictions should be interpreted as model-based extrapolations rather than definitive forecasts. Uncertainty generally increases as the prediction horizon becomes longer, and long-horizon predictions can be

sensitive to model assumptions about trend, seasonality, and autoregressive dependence.

The examples above illustrate the basic univariate workflow for fitting, diagnosing, visualizing, and making short-term predictions from a `tempssm` model. The detailed manual extends this workflow to exogenous variables, additional model diagnostics, and time-series cross-validation:

https://github.com/akihirao/tempssm/blob/main/tools/manual/tempssm_manual.pdf